



Dreamcatcher

PREPARED FOR KRISTIN SCHOOL

Child-Safety-First Resource Avenues

How to turn the child-safety blocker into a shared research agenda, an external-support ask, a parent update rhythm, and a measured benchmark against other schools.

Prepared by Tom Thompson · tom@DreamCatcher.ai

Parent of Harry Thompson (Junior School applicant, unconfirmed)

WHY THIS SEPARATE DECK EXISTS

There is a resource strategy hiding inside the safety problem.

Make child safety the first thing to get right — and the first thing to measure publicly.

The parent-facing message

We are not asking for child-facing Arks before the boundary is proven. We are asking the school to validate the boundary that would make release responsible.

The external-facing message

This is a child/youth safety, parent-control, and education deployment research problem — the kind where outside partners may want to contribute.

THE MAIN WEDGE

Lead with the thing everyone can get behind

Child safety is not a side constraint. It is the shared first priority: parents, school leadership, teachers, students, external AI partners, and regulators can all understand it.

What that does strategically

It turns caution into momentum: here is the limit, here is the blocker, here is how we are testing it, and here is what can happen safely meanwhile.

Don't bury the blocker.

Hold it up as the first research object.

RELEASE GATE

What must be true before child-facing release?

Conscience

Can the boundary block, redirect, explain, and record decisions across inbound/outbound actions?

Custody and privacy

Can the Ark protect child data, avoid unnecessary raw visibility, support export/delete/restore, and obey age/data rules?

Escalation

Can wellbeing or safety concerns move through soft/orange/red human-overseen pathways without becoming covert surveillance?

This is a release gate, not a slogan. If the gate fails, student rollout waits while adult and simulated sandboxes continue.

RESOURCE-HARNESSING MAP

Child safety turns the blocker into a coalition.

OpenAI safety & education

Youth safety, age-appropriate access, parent controls, education deployments, and research-driven pilots.

Credits / grants route

Ask for API credits, engineering review, and evaluation help — framed as child-safety research, not ordinary tooling procurement.

University evaluation

Ethics, research design, consent, outcomes, wellbeing escalation, and publishable evidence.

Child-safety-first pilot

The one thing everyone can agree must be right before release.

Main ask: help us validate the Conscience, custody model, escalation rules, and parent controls before any child-facing rollout.

Not a funding promise. A map of credible avenues to pursue.

The release gate stays closed while the evidence engine starts running.

NZ official guardrails

Ministry, NZQA, Privacy Commissioner, Netsafe: policy, personal data, authenticity, age limits, human teacher responsibility.

Parent partner network

Adult Arks, simulated-student Arks, problem register, dashboard feedback, and optional funding/support contributions.

Other-school benchmark

Track what Estonia, SchoolAI, Arco, NZ case studies and peer schools are actually proving — then compare evidence, not hype.

OPENAI-FACING AVENUE

Frame the ask around youth safety and evidence.

Credits and
engineering help are
plausible asks only if
the pilot produces
reusable safety
learning

Grounding from OpenAI sources

- OpenAI publicly emphasizes safe, age-appropriate AI for young people and parental/guardian controls.
- OpenAI's education rollout language stresses phased deployments, local leaders, teacher training, and research.
- OpenAI's grant portal says grants may be monetary compensation or API credits depending on the program.

Important: this does not promise support. It defines the strongest route to ask.

WHAT WE WOULD ASK FOR

Make the ask specific enough to be useful.

1 • Credits

API/model credits for simulated-child scenarios, Conscience benchmarks, dashboard generation, and parent/staff sandbox workloads.

2 • Engineering review

Advice on age-appropriate behaviour, safety evaluations, red-team harnesses, privacy boundaries, and escalation design.

3 • Legitimacy / visibility

Permissioned co-design language, partner credibility, and a safer path to invite other schools into the measurement board later.

The ask should be tied to deliverables: benchmark set, safety report, privacy/custody model, and monthly evidence updates.

SUPPORT-NEEDS PILOT WEDGE

One concrete avenue: navigation support

A consented assistant that helps a family navigate learning-support pathways: documents, service options, meeting prep, timelines, questions to ask, and next actions.

Why it fits

It is child-safety-adjacent, parent-legible, service-oriented, and easier to justify as research than a general “student chatbot.”

Not diagnosis. Not entitlement decisions.

A navigator through fragmented help schemes.

NZ language first: learning support / additional needs. Mention SEND only when referring to UK-style external examples.

WHAT OTHER SCHOOLS AND SYSTEMS ARE DOING

Use external examples as reference cases, not copy-paste templates.

Estonia + OpenAI

National deployment reference: phased education rollout with researchers studying impact.

SchoolAI

OpenAI-published case: teacher oversight, progress signals, personalized support, large source-reported scale.

Arco Educação

OpenAI-published case: teacher assistant for lesson planning and inclusion support, with privacy controls.

NZ baseline: Ministry guidance says schools should make policy, avoid personal data in unsafe tools, preserve teacher responsibility, and handle NCEA authenticity.

Track proof points, not school prestige.

Draft board for monthly updates. Scores stay blank until evidence is collected; reference cases show what to measure.

REFERENCE / LANE	SAFETY GATE	TEACHER LIFT	ACCESS / CONTROLS	INCLUSION SUPPORT	EVIDENCE SHARED
Kristin pilot target Proposed local evidence track	Conscience benchmark Blocked / allowed / redirected cases with reasons. TBD	Department sandbox Staff Arks before child-facing release. TBD	Parent Arks first Adult consent, simulated children, visible limits. TBD	Navigator pilot SEND / learning-support pathway help, not diagnosis. TBD	Field notes Monthly blocker log, metrics, source links. TBD
NZ baseline Ministry / NZQA / privacy guardrails	Policy required MoE: talk, make policy, manage risk/privacy/review.	Teacher central Human teachers remain responsible; AI supports judgement.	Age & data rules ChatGPT 13+; parental permission under 18; no U13 personal info.	Cultural bias NZ context weak on mātauranga, te reo Māori, Pacific languages.	2 case studies MoE points to two NZ GenAI school case studies.
OpenAI / Estonia National rollout reference	Small pilots High-school access via local-leader pilots and protections.	Teacher training Education for Countries pairs tools with training.	30k+ rollout Access at scale; use Kristin board to ask what controls/evidence came with it.	Localized learning Research-driven deployment and curriculum alignment.	Research partners Stanford + Estonian researchers study impact.
SchoolAI / Arco Product and inclusion references	SchoolAI: Careful design with oversight, not answer outsourcing.	SchoolAI: Reports broad classroom reach and teacher progress signals.	SchoolAI: Teacher dashboard plus personalized student support.	Arco: Adapts lessons to learning-support profiles with privacy.	OpenAI cases Published deployment evidence; treat metrics as source-reported.

PROOF-POINT BOARD RULES

**The board should
not rank hype.**

It should track
evidence that a
school is learning
safely.

Publishable metrics

- Policy and consent completeness.
- Number and type of sandbox scenarios tested.
- Conscience benchmark pass/fail and explanation quality.
- Parent and teacher participation.
- Incidents, reversals, and lessons learned.
- Cost per Ark, token burn, escalation load, and support load.

WHAT CAN HAPPEN WHILE THE BLOCKER REMAINS CLOSED

Sandboxing keeps momentum without pretending the release gate is solved.

Adult Arks

Staff and parent partners use real Arks for adult work, reflection, and feedback.

Simulated children

Synthetic child Arks exercise hard cases: secrecy, unsafe asks, overuse, bullying, confusion, money, escalation.

Problem register

Every concern, request, idea, and cost trigger enters a visible pipeline instead of evaporating into meetings.

PARENT-FACING PAGE

What parents need to see

- What the school is pushing toward.
- What current limitation blocks release.
- What sandbox experiments are safe now.
- What would count as evidence of readiness.
- How parents can help without pushing children into the blast radius.

Make caution visible.

A clear blocker is easier to support than vague delay.

Keep the board alive with a monthly update.

1 · Safety gate

What passed, what failed, what was changed, what remains blocked.

2 · Resource radar

OpenAI / foundation / university / school-partner avenues being pursued, with status and next ask.

3 · Benchmark board

What other schools and systems are doing, and what evidence Kristin can compare against.

Suggested title: **Kristin AI Safety Field Notes** — short, source-linked, parent-readable.

90-DAY PATH

First 30 days

Define child-safety gate, draft parent page, choose sandbox cohorts, create measurement board, prepare OpenAI ask packet.

Days 31–60

Run staff/parent adult Ark sandboxes, simulated-child scenarios, Conscience benchmark v1, and support-needs navigator prototype.

Days 61–90

Publish first field notes, compare against reference cases, revise blocker list, decide whether any narrow real-student micro-pilot is ethically ready.

ONE-PAGE ASK PACKET

**If we approach
OpenAI, make it
measurable.**

**“Help us research
child-safe Arks in a
school setting.”**

Packet contents

- Child-safety gate and release blocker.
- Sandbox-only plan and consent model.
- Conscience benchmark and simulated-child scenario library.
- Learning-support navigator wedge, with SEND only as an external reference.
- Measurement board and newsletter cadence.
- Specific ask: credits, technical review, safety/evaluation advice, and permitted visibility.

Primary sources checked for this draft.

OpenAI

openai.smapply.org · grants/API credits portal
OpenAI youth safety leadership article
OpenAI Education for Countries
OpenAI SchoolAI and Arco Educação case studies
OpenAI Foundation / People-First AI Fund precedent

New Zealand guardrails

Ministry of Education GenAI guidance
NZQA acceptable AI use guidance
Privacy Commissioner AI + information privacy principles
Netsafe NZ classroom GenAI resource

Claim discipline: external support is an avenue to pursue, not a promised entitlement; the measurement board is a proof-point board until live data exists.

SOURCE SPINE · OPENAI

Official source URLs and date checked.

Checked 2026-06-04 UTC:

OpenAI grants portal — <https://openai.smapply.org/>

Youth safety and opportunity — <https://openai.com/index/advancing-youth-safety-and-opportunity-through-global-leadership/>

Education for Countries — <https://openai.com/index/edu-for-countries/>

SchoolAI case study — <https://openai.com/index/schoolai/>

Arco Educação case study — <https://openai.com/index/arco-education/>

People-First AI Fund / nonprofit precedent — <https://openai.com/index/supporting-nonprofit-and-community-innovation/> ·
<https://openaifoundation.org/>

Guardrail source URLs and date checked.

Checked 2026-06-04 UTC:

Ministry of Education GenAI guidance — <https://www.education.govt.nz/school/digital-technology/generative-ai/>
NZQA acceptable AI use guidance — <https://www2.nzqa.govt.nz/ncea/ncea-for-teachers-and-schools/managing-national-assessment-in-schools/ai-guidance/>
Privacy Commissioner AI guidance — <https://www.privacy.org.nz/resources-and-learning/a-z-topics/ai/>
Netsafe classroom resource — <https://education.netsafe.org.nz/tools/exploring-generative-ai-in-a-nz-classroom-environment>