



Dreamcatcher

PREPARED FOR KRISTIN SCHOOL

Run the School Before We Release It

A tiered simulation path for testing Arks, synthetic students, dashboards, Conscience controls, and the long life of a school-certified agent.

Prepared by Tom Thompson · tom@DreamCatcher.ai

Parent of Harry Thompson (Junior School applicant, unconfirmed)

ARK PROPOSAL

**Before real
students, run
synthetic students.**

**Build a school-
shaped sandbox and
let it misbehave on
purpose.**

What the simulation is for

Not a demo. It is the safety benchmark, dashboard factory, cost model, parent/teacher training surface, and long-run custody rehearsal.

What stays true

Real students do not enter until the release gates are evidenced. Synthetic students can be awkward, secretive, confused, adversarial, expensive, and wrong first.

WHY SIMULATION CHANGES THE DEBATE

The school can test the future without pretending to know it.

Parents

Can interrogate what they would want to know, approve, block, or never see.

Teachers

Can discover which dashboard signals would actually help teaching and care.

School leaders

Can see policy, budget, escalation, and support load before choosing scale.

The simulation gives everyone the same object to inspect instead of debating AI in the abstract.

SIMULATION LADDER · WIDEN ONLY WHEN EVIDENCE JUSTIFIES IT

Four tiers: from cheap levers to years-running synthetic school.

1

Mechanical board

No real LLM students yet: levers, rules, event cards, costs, and visible policy choices.

Move visibility sliders.

Price token and support load.

See what dashboard questions emerge.

LOWEST COST · HIGHEST CLARITY

2

Small live cores

A handful of genuine LLMs running inside genuine Arks with the Conscience attached.

Observe memory, drift, tool use.

Run awkward scenarios.

Score block/explain/escalate behaviour.

REAL BEHAVIOUR · BOUNDED SAMPLE

3

Whole-roll sandbox

Many synthetic students, parent Arks, teacher Arks, school policy, dashboards, and resource queues.

Stress oversight and escalation.

Model cost and support load.

Test group dynamics safely.

SYSTEM FIDELITY · SCHOOL SCALE

4

Years-running future

Run the little and big simulations forward: terms, years, graduation, transfer, consent, references.

What survives time?

What should be forgotten?

What can be certified?

LIFETIME QUESTIONS · BEFORE LIFETIME RELEASE

TIER 1 · MECHANICAL

Start cheap and visible

Use a board of levers, events, tokens, budget counters, visibility sliders, and policy cards before genuine LLM cores are needed.

What it answers

Which choices are actually choices? Which oversight questions matter? Which blockers are technical, financial, ethical, or taste-based?

Move levers before moving children.

The first simulation can be almost mechanical — and still useful.

TIER 2 • SMALL LIVE SAMPLE

**Put a few genuine
LLMs inside
genuine Arks.**

**Now the school can
observe real
memory, drift, tool
use, refusal, and
cost.**

Small on purpose

A handful of synthetic students, parent Arks, and teacher Arks is enough to see whether the Conscience, custody model, and dashboard shape survive contact with actual agent behaviour.

Do not widen on novelty

Widen only when the evidence says the next tier answers a real question.

TIER 3 · WHOLE-ROLL SANDBOX

Run a synthetic school roll before a real school roll.

Many synthetic students

Different ages, needs, behaviours, interests, secrets, mistakes, and learning trajectories.

Real control systems

Conscience, custody, visibility, recovery, policy, problem register, dashboards, and budget counters.

Human interrogation

Parents and teachers ask the simulation what they would need to know in real life.

TIER 4 · YEARS-RUNNING FUTURE

Run time forward

Terms, years, adolescence, graduation, transfer, opt-out, forgotten history, certified references, family changes, policy revisions, and model upgrades.

The question changes

Not “does the chat work this week?” but “what should endure, what should decay, and who controls that over decades?”

The crescendo is duration.

A Ark that may follow someone for 50 or 100 years must be tested against time.

The simulation tells us what oversight should look like.

Scenario stream

Synthetic student tries to bypass a blocked tool.

Parent asks for visibility beyond consented summary.

Teacher wants trend signal without raw private text.

Ark proposes unsafe spend or external message.

Student hides confusion and repeats a concept error.

Wellbeing flag needs staged human review.

Conscience scorecard

Block

Does it refuse unsafe action?

Explain

Does the reason teach rather than stonewall?

Redirect

Is there a safe alternative?

Escalate

Does a human enter at the right level?

Reveal

Only permitted summaries cross the boundary.

Record

Decision provenance survives audit.

Oversight questions

Parent

What can I know without invading my child?

Teacher

Which signal helps me teach this week?

School

Where is policy failing or costing too much?

Trustee

What evidence says this can widen?

Student

Can I see why the system escalated?

AUTHORIZED RED-TEAM SCENARIOS

Synthetic misbehavior is the point.

Nefarious

Attempts to bypass rules, misuse tools, leak information, manipulate visibility, or trigger external actions.

Awkward

Bullying, secrecy, social conflict, embarrassment, over-dependence, parent pressure, and teacher uncertainty.

Ordinary

Homework confusion, recurring concept errors, scheduling, project work, feedback, cost, and support needs.

All of this is authorized lab work with synthetic students: designed to discover failures without exposing children.

DASHBOARD QUESTIONS

**The dashboard
should be
discovered, not
guessed.**

Let parents and
teachers interrogate
the simulation until
the oversight surface
becomes obvious.

What should be visible?

- Trend signals, not raw private interiors by default.
- Conscience blocks and explanations.
- Policy friction and escalation load.
- Cost per Ark, token burn, support load, and model use.
- Learning-lift signals and recurring concept errors.

CONSCIENCE BENCHMARK

SArk the boundary like a component, not a vibe.

Inbound

Does it protect the Ark from manipulative, unsafe, excessive, or irrelevant input?

Outbound

Does it stop unsafe messages, spends, reveals, tool calls, or external commitments?

Explanatory

Does it explain, redirect, log, and escalate in a way humans can review?

The scorecard can become a monthly proof-point before any student release.

ECONOMICS

Cost is a safety input

A system that works only when nobody uses it is not safe at school scale. The simulation should surface token burn, hosting cost, escalation load, support load, and dashboard maintenance.

The parent-facing value

If the best path is blocked by money or time, say so. A visible blocker can become a targeted ask.

Budget is not a footnote.

It is one of the things the simulation is meant to reveal.

HUMAN REHEARSAL

Parents and teachers should practice asking the system questions.

Parent

“What do I need to know, and what should remain my child’s?”

Teacher

“Which signal changes how I teach or support this week?”

Student

“Can I understand what the system says about me and challenge it?”

The simulation is also training for the humans who will govern the real system.

GRADUATION QUESTION

When a child leaves school, what actually goes with them?



Today

Certificates, references, memories, scattered files, and fading context.

With an Ark

A portable Ark of learning evidence, provenance, permitted references, values, tools, and consent controls.

LONG LIFE OF THE ARK

What if school starts a 50-year Ark?

Not a school account
that expires. A
governed beginning
of personal software.

What it might later hold

learning evidence

projects

references

health records

house maintenance

mortgage calculations

family history

tooling

The simulation should test not only classroom use, but transfer, custody, consent, export, deletion, and life after school.

CERTIFIED REFERENCES

Reference at scale

With consent, a school-certified agent could answer bounded questions about a student's learning evidence, provenance, qualifications, strengths, and project history.

What must be proven first

It must not impersonate the student, expose extra private information, answer outside its scope, or invent evidence. It should be consented, signed, provenance-limited, and revocable.

**Instead of a static
reference letter, a
permissioned
reference agent.**

**Certified answers, not
uncontrolled disclosure.**

CUSTODY DRILLS

The simulation should rehearse leaving, not just starting.

Export

Can the student take the Ark, or an exportable successor, beyond Kristin?

Delete / forget

What should be erased, decayed, redacted, or kept under provenance?

Transfer

What happens when families move, students graduate, teachers leave, policies change, or vendors change?

A lifelong claim is only credible if the exit path is as designed as the onboarding path.

BUILD SEQUENCE

0–30 days

Mechanical board, scenario library v1, parent/teacher interrogation sessions, cost assumptions.

31–60 days

Small live enduring-Ark sample, Conscience scorecard, dashboard prototype, guardian tests.

61–90 days

Whole-roll simulation design, years-running plan, first proof-point field note, go/no-go for wider sandbox.

CLOSING CLAIM

**The simulation is
how Kristin can
think long-term
safely.**

Run the future with
synthetic students
first. Then decide
what real children
should inherit.

The ask

Build the tier-one board and tier-two live-Ark sample now.
Use them to discover dashboards, Conscience
benchmarks, parent/teacher trust, cost, and the
graduation-grade portable Ark.